Building and Evaluating Predictive Models – Part 2

Abhishek Kumar
AUTHOR

@meabhishekkumar

Data Science Project Cycle

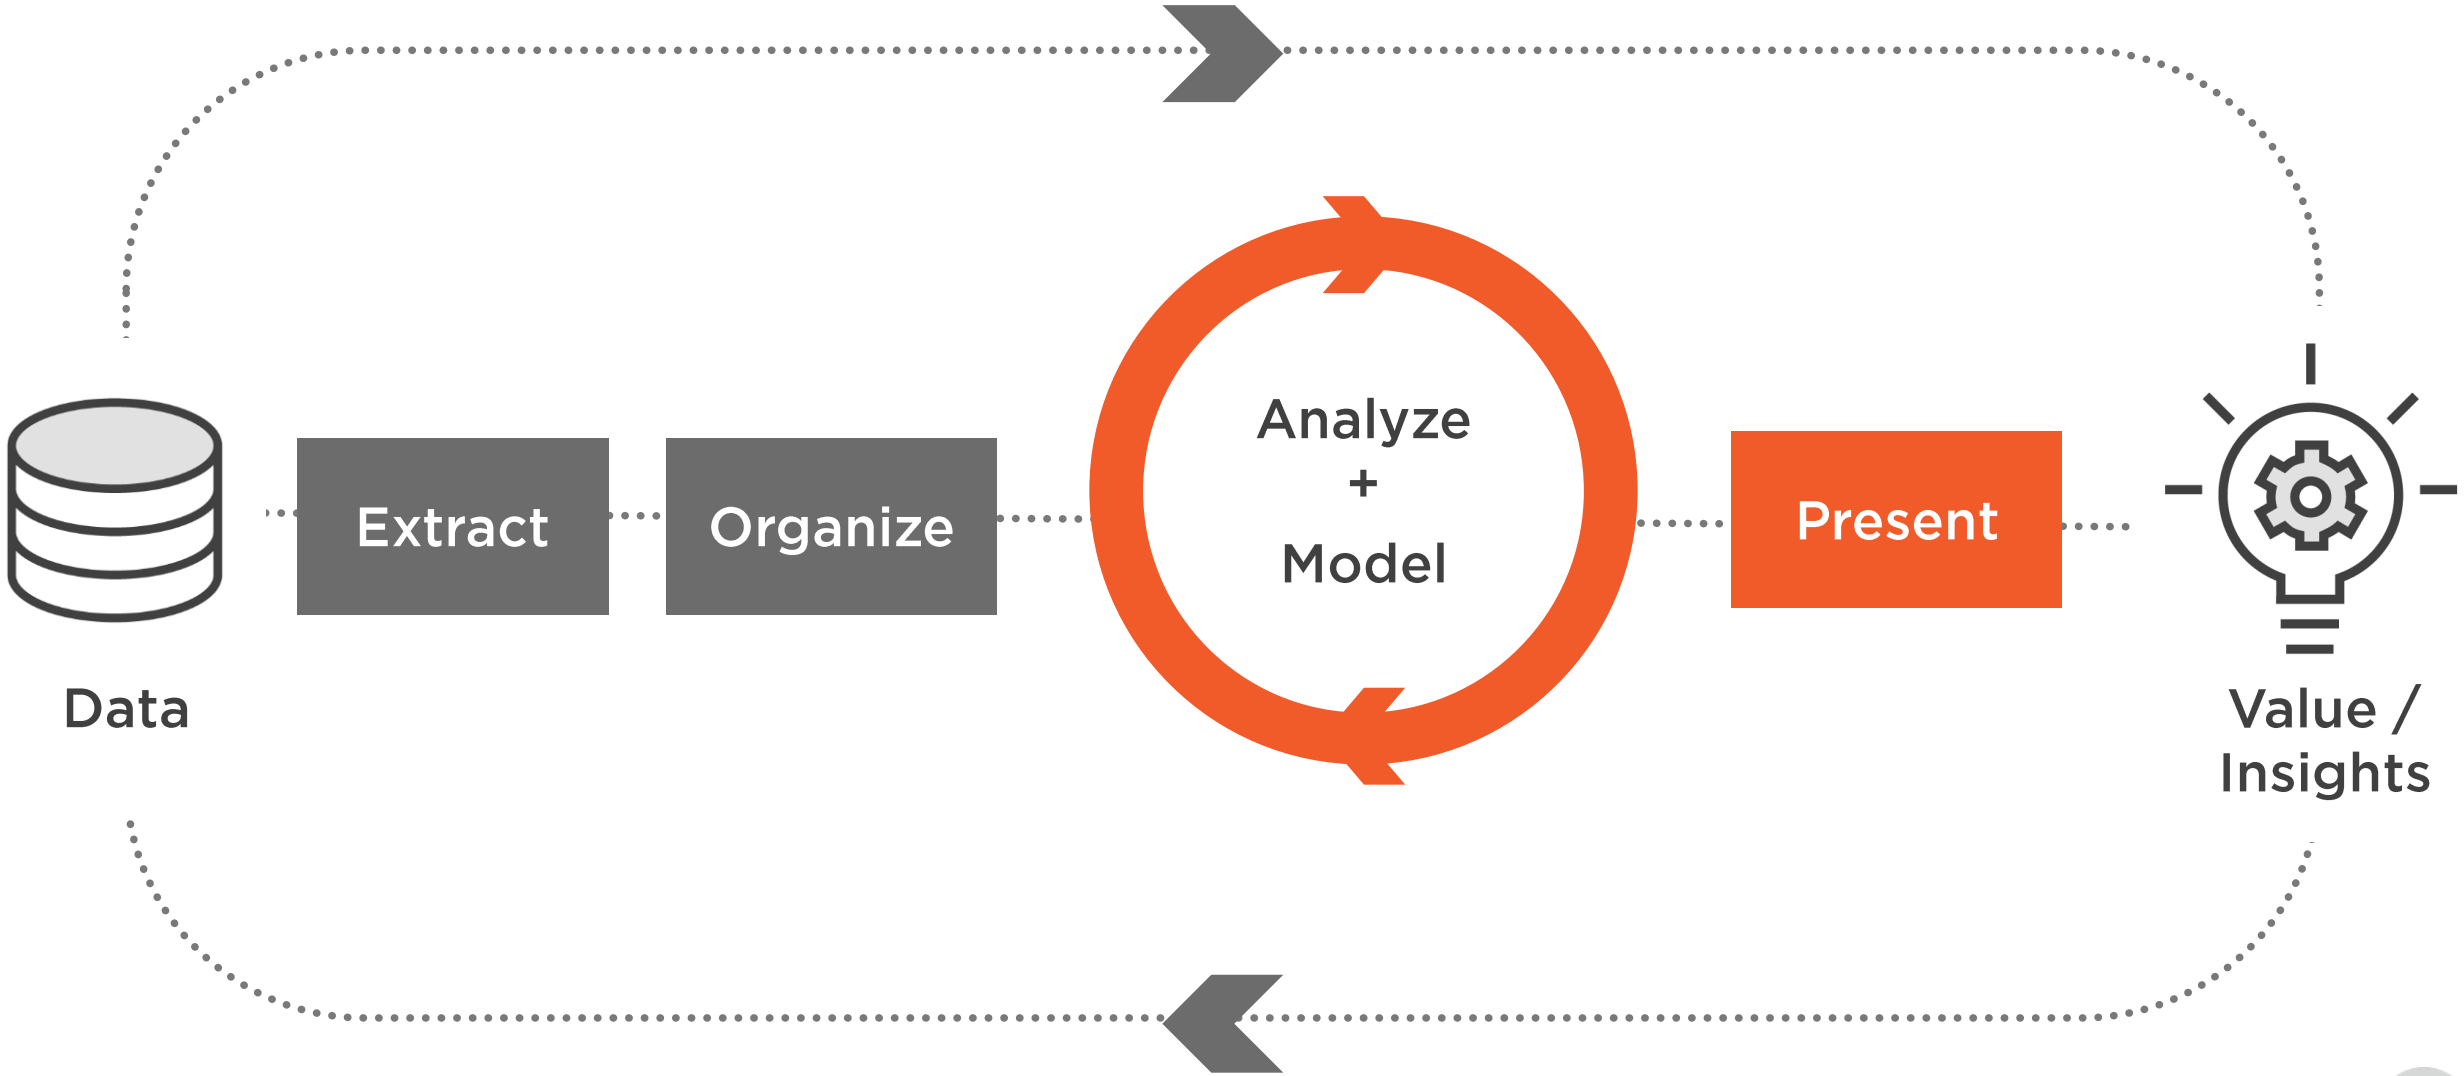

Data Science Project Cycle

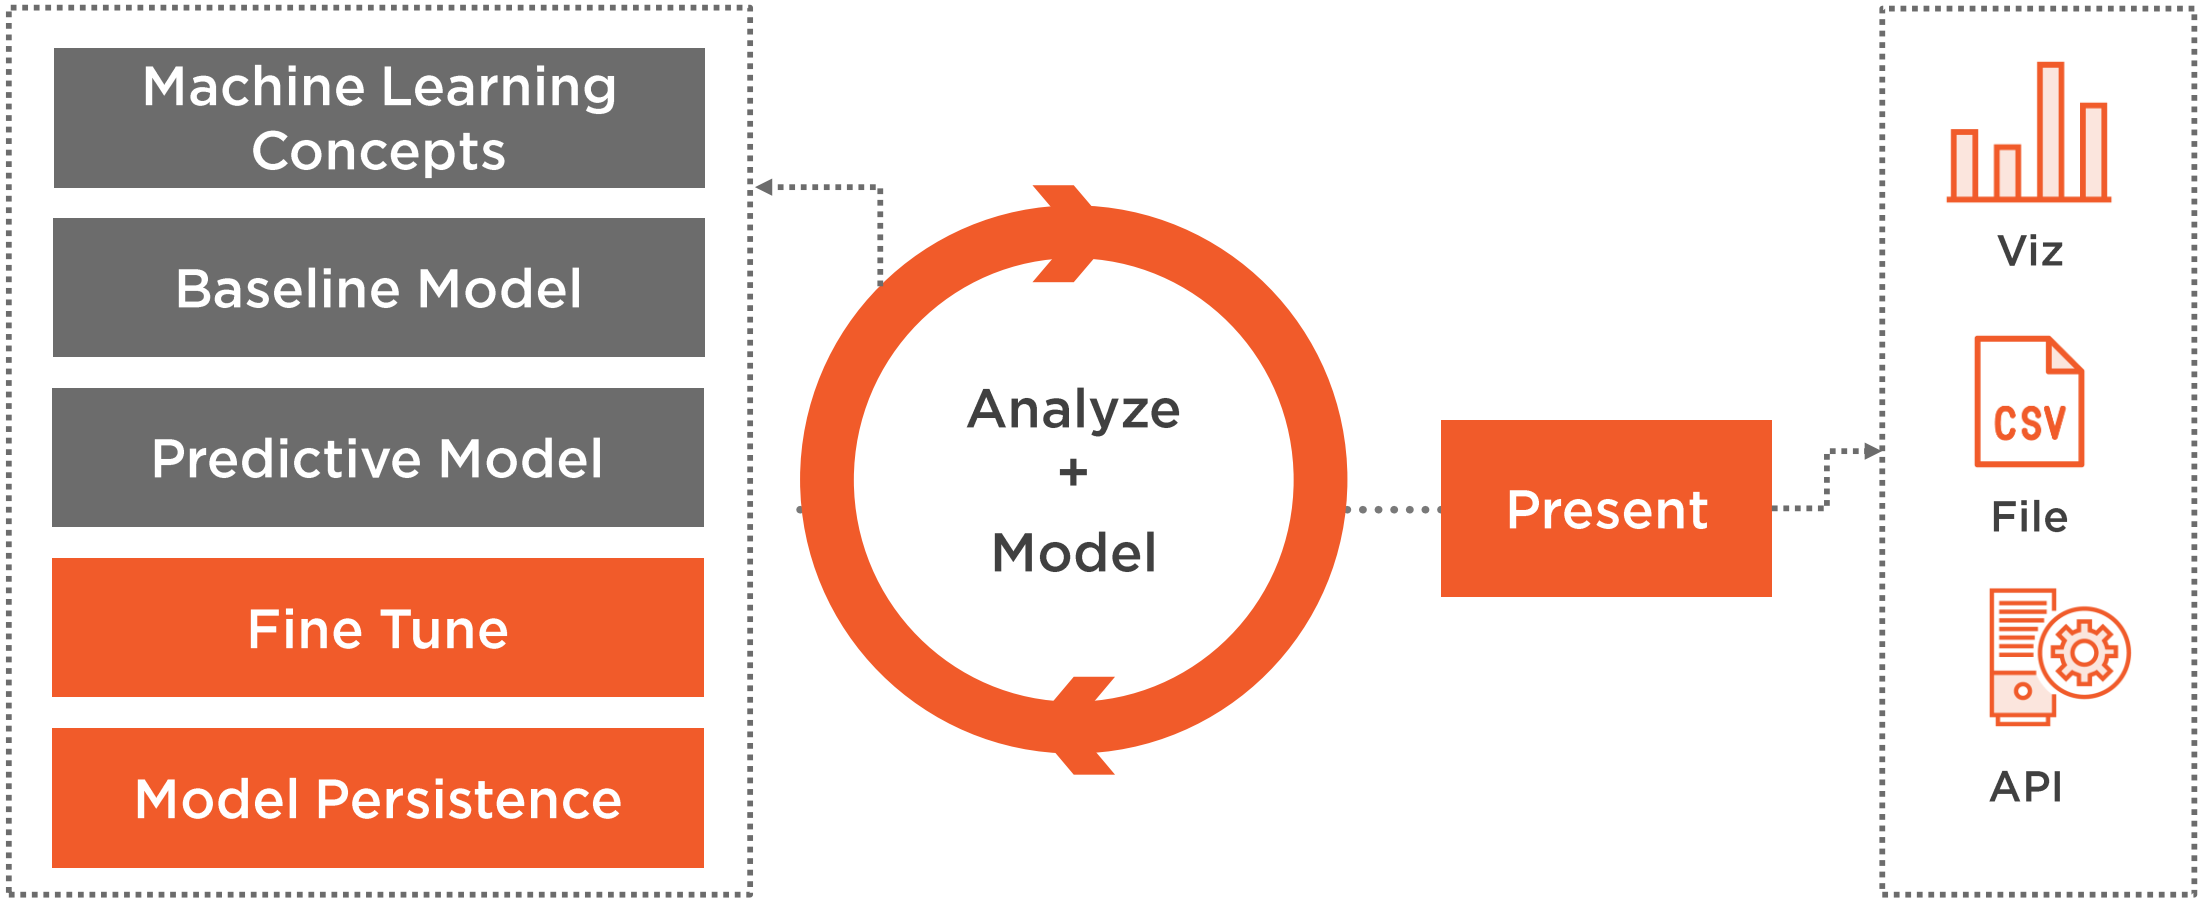

Overview (Contents)

Model tuning

- Underfitting vs overfitting
- Regularization
- Hyperparameter tuning
- Cross validation

Feature Engineering

- Feature normalization

Model persistence

API

Overview (Tools)

Python

- Numpy
- Pandas
- Scikit-Learn
- Pickle
- Flask

Underfitting vs. Overfitting

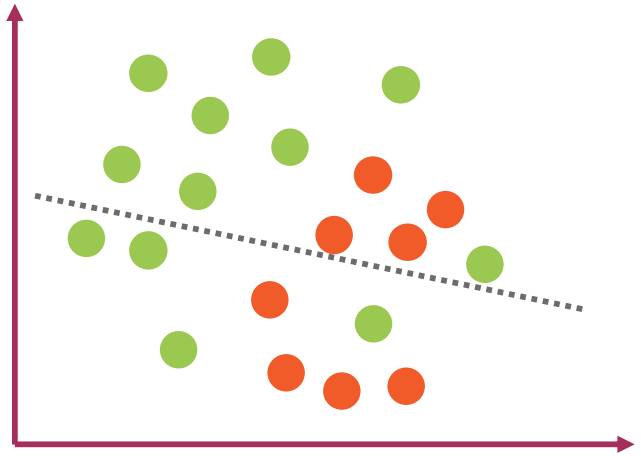

Underfitting

- Can’t learn the pattern in the training data

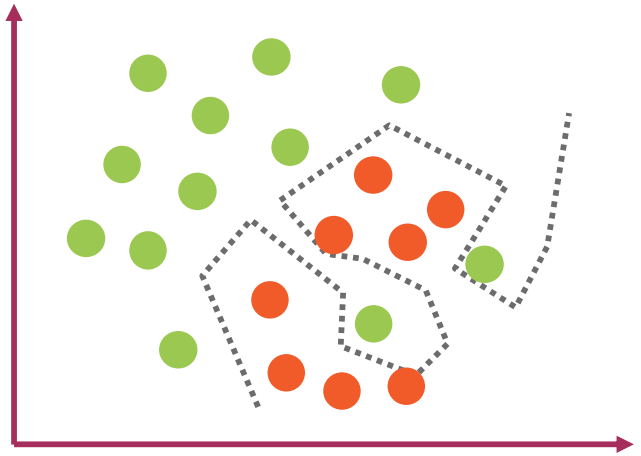

Overfitting

- Memorize training data
- Poor generalization

Regularization

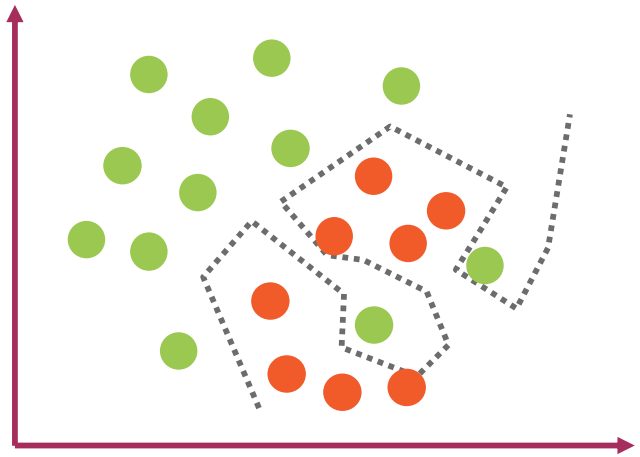

Overfitting

Regularization

Reduce model complexity

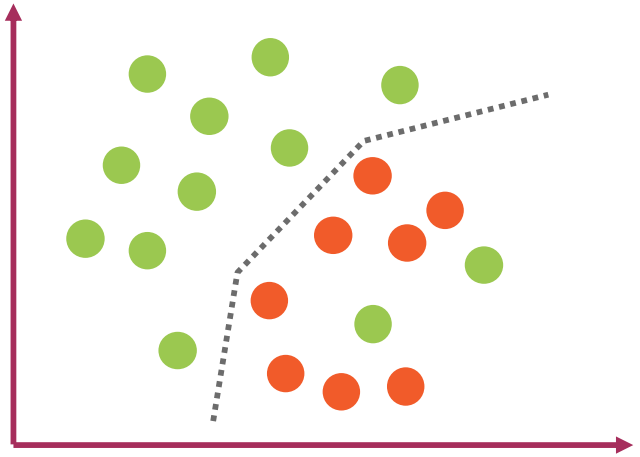

Balanced

Regularization

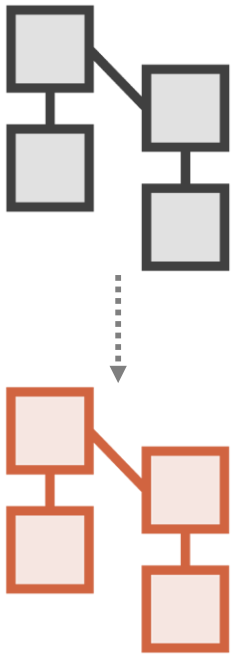

Create model

```
model = LogisticRegression(random_state=0)
```

Train model

```
model.fit(X_train, y_train)
```

Score

```
model.score(X_test, y_test)
```

Coefficients (Model parameters)

```
model.coef_
```

Regularization

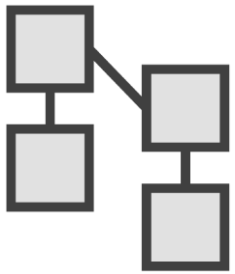

Create model

model = LogisticRegression(random_state=0)

Regularization parameter

Large → **Overfit**

Increase model complexity

Small → **Underfit**

Decrease model complexity

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
          intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
          penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
          verbose=0, warm_start=False)
```

Regularization

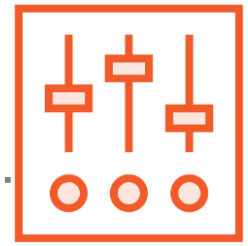

Hyperparameter

Hyperparameter Optimization

Regularization parameter

Large

Overfit

Increase model complexity

Small

Underfit

Reduce model complexity

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
    verbose=0, warm_start=False)
```

L1

L2

Hyperparameter Optimization : GridSearch

Model (A)

a1	a2	a3

Model (A, B)

	a1	a2	a3
b1			
b2			
b3			
b4			

Create grid with different combinations → Evaluate each combination → Select combination with best model performance

Train-test Split

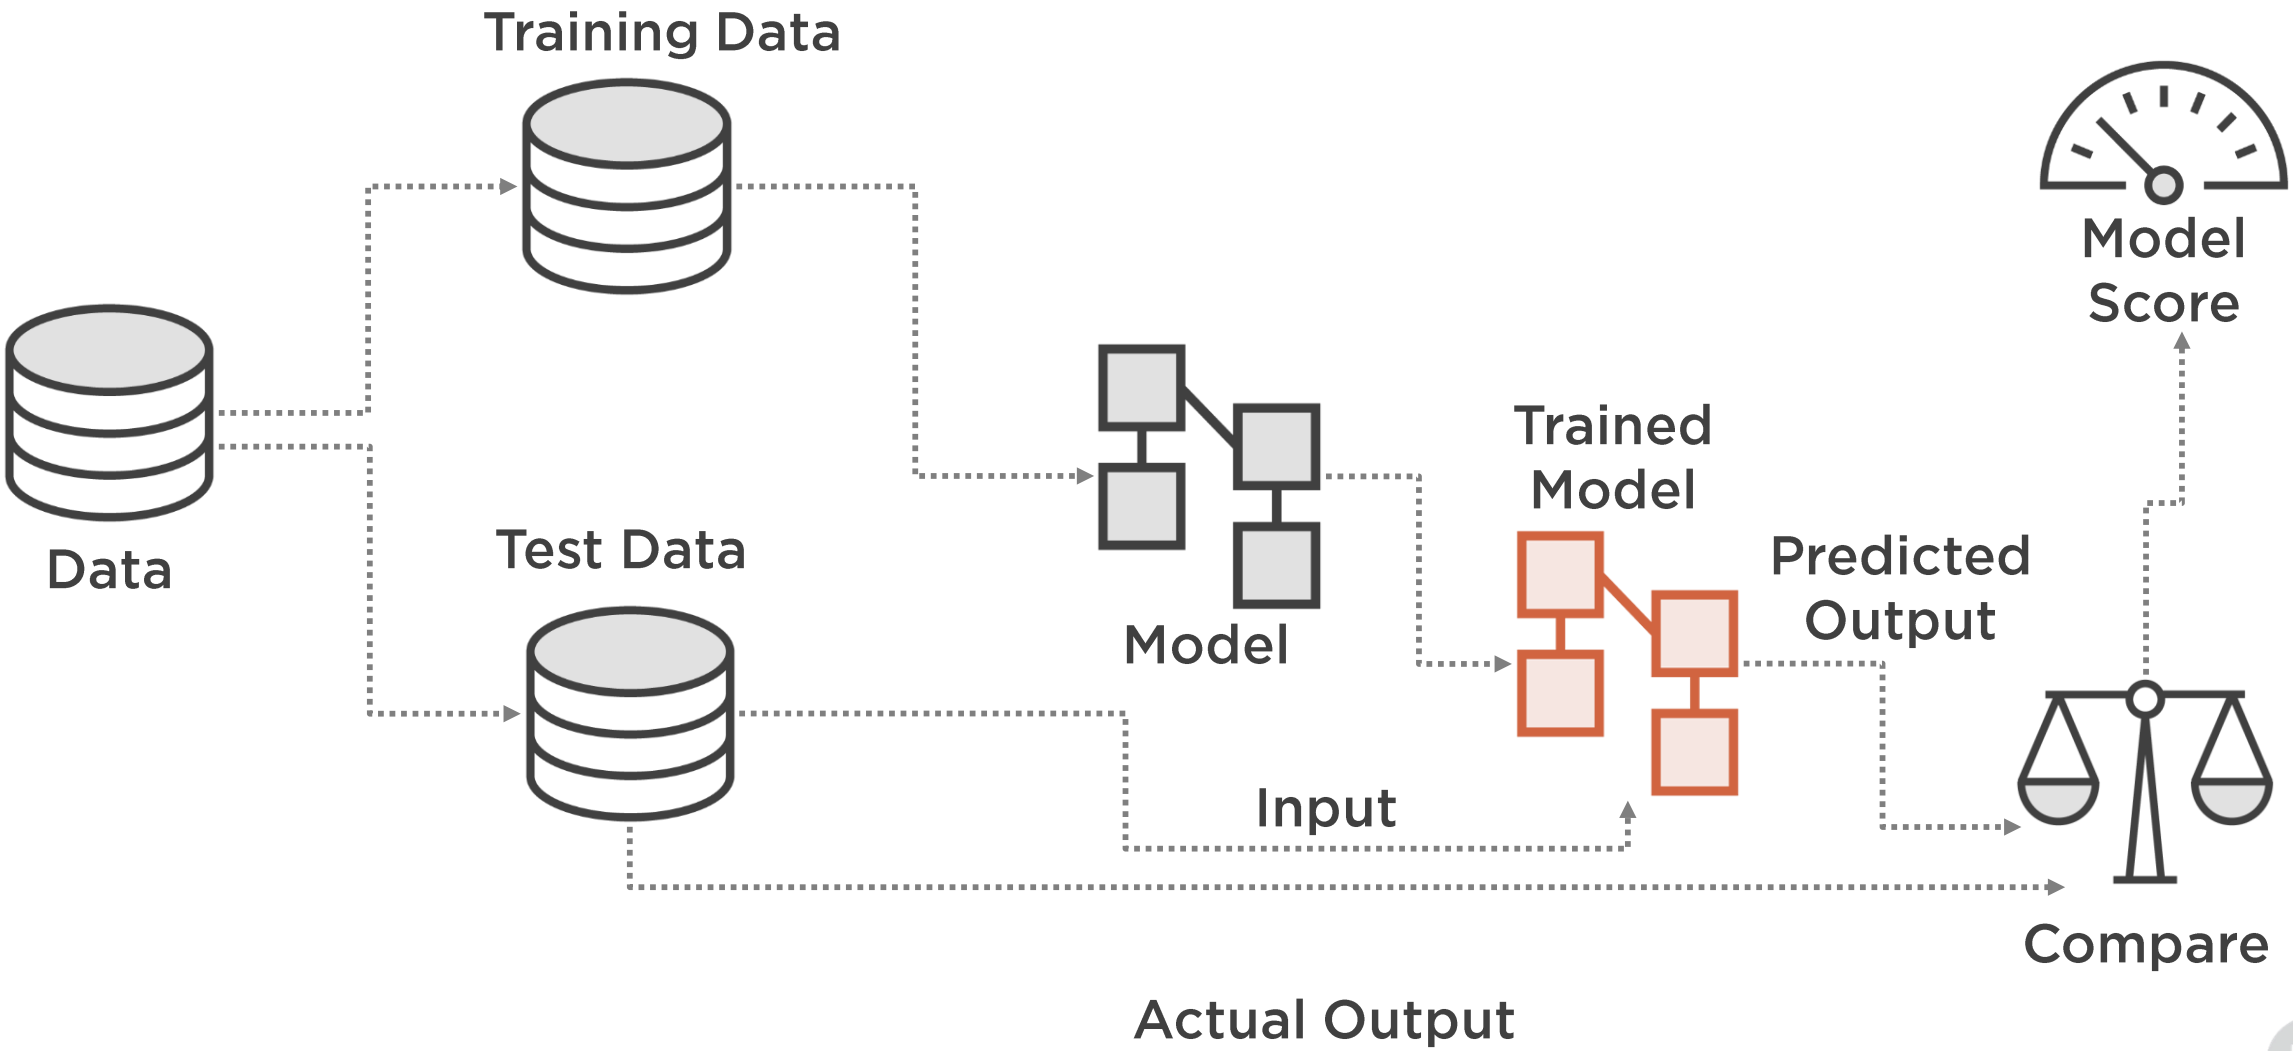

Cross-validation

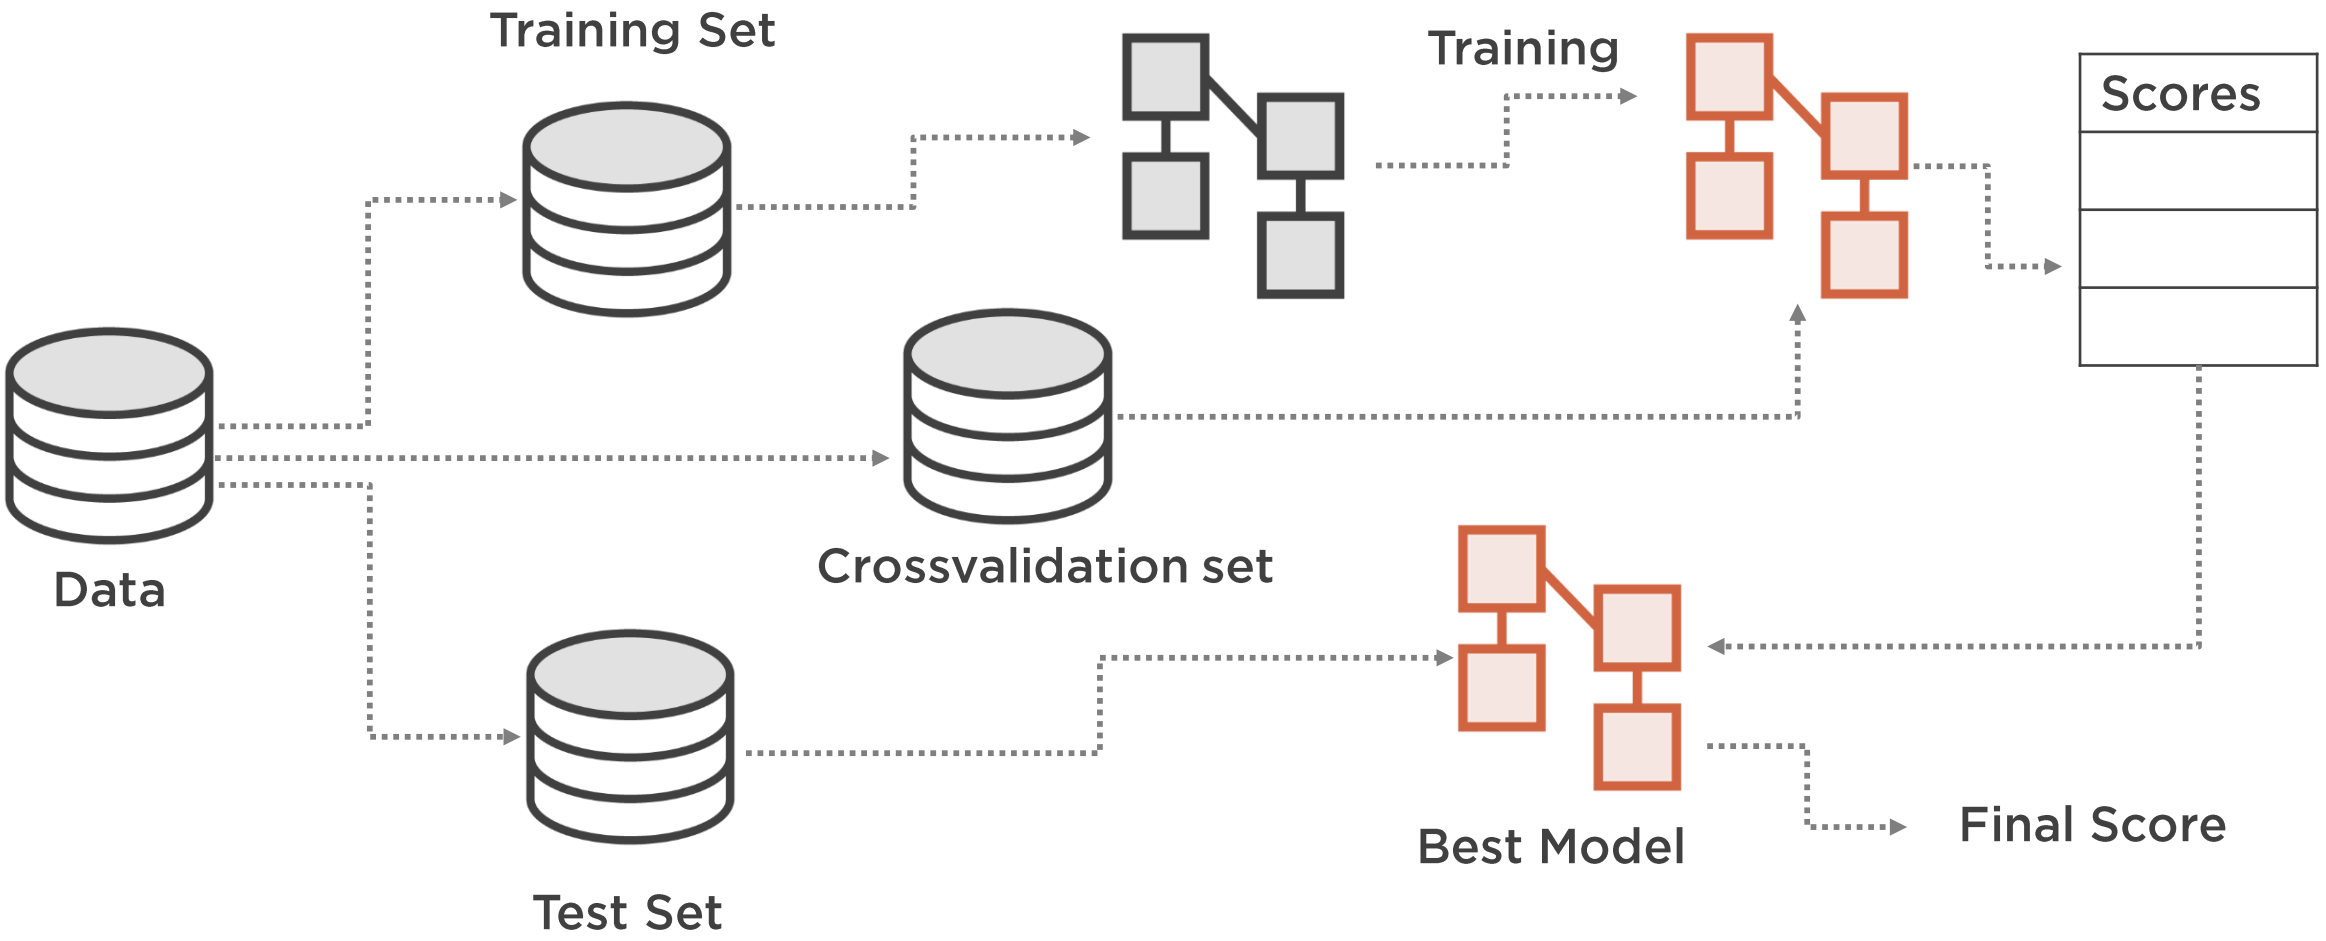

K-Fold Cross-validation

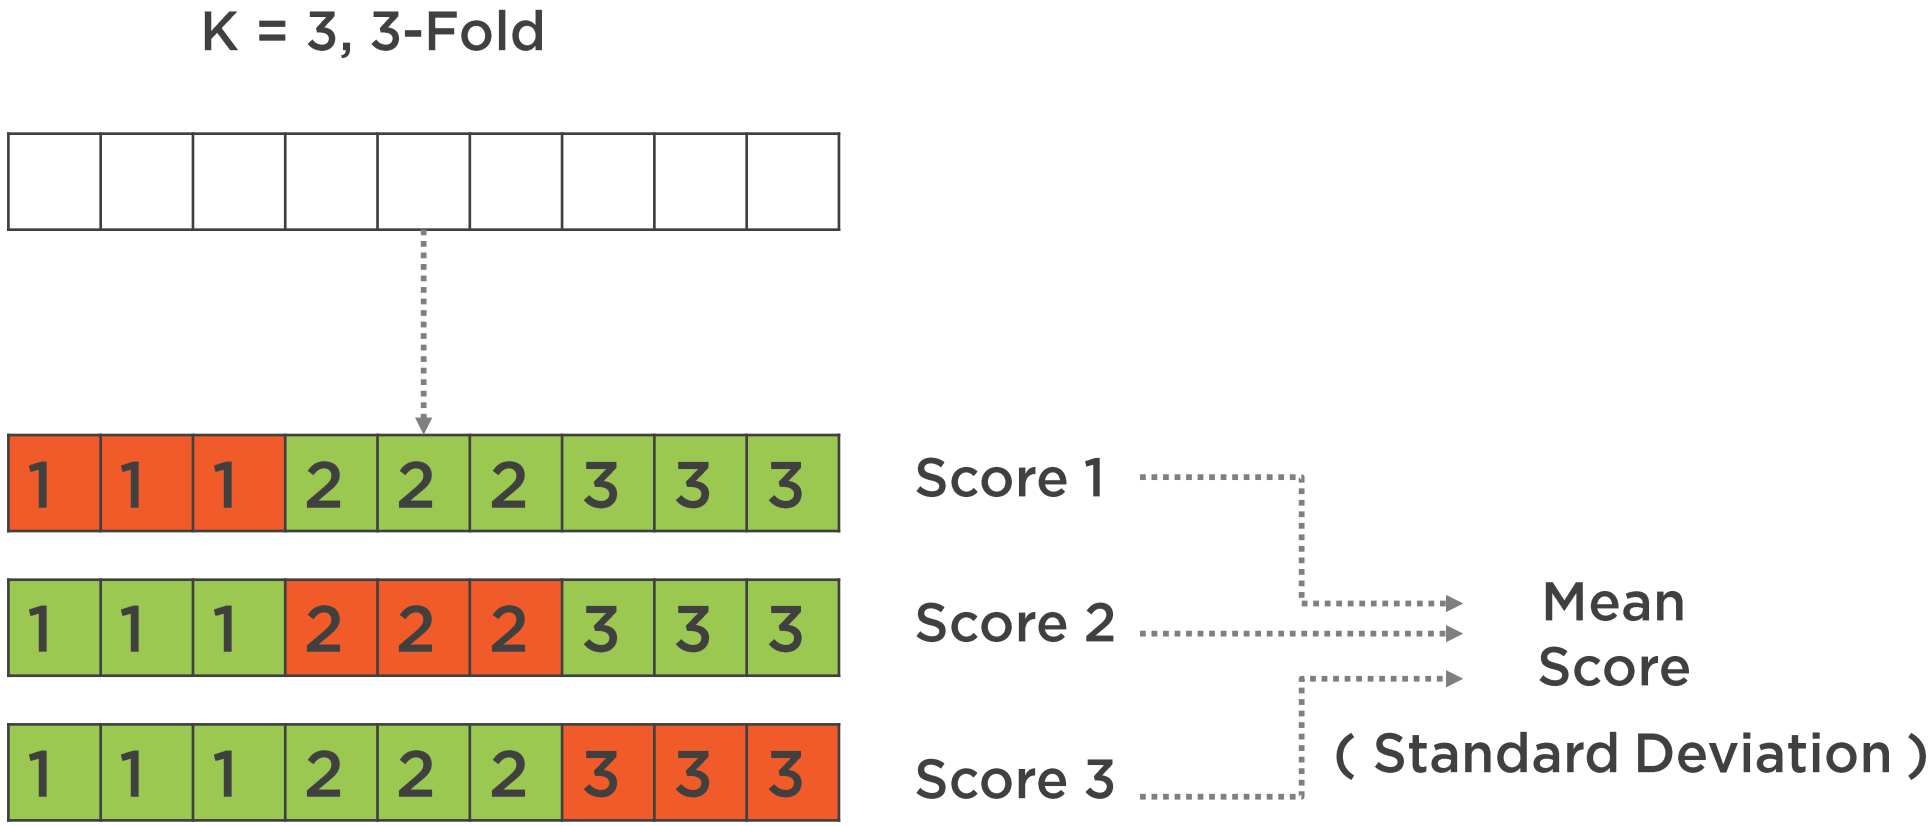

Demo

Hyperparameter optimization using GridSearchCV

Demo

Making third Kaggle submission

Feature Normalization

Age	Fare	FamilySize
..	..	..
..	..	..
..	..	..

0.4 to 80	0 to 512	1 - 11	
0 to 1	0 to 1	0 to 1	Scale Type 1
-1 to 1	-1 to 1	-1 to 1	Scale Type 2

Feature Standardization

0.4 to 80	0 to 512	1 - 11
Age	Fare	FamilySize
..	..	..
..	..	..
..	..	..

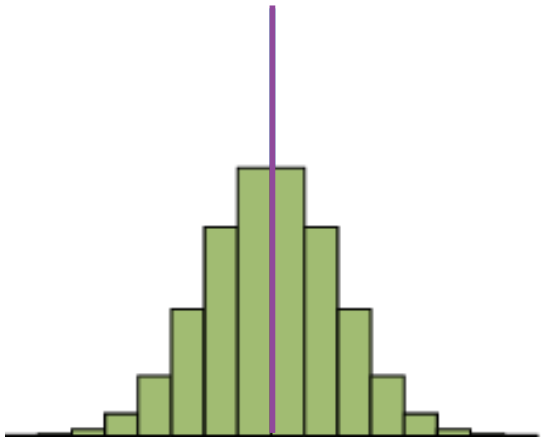

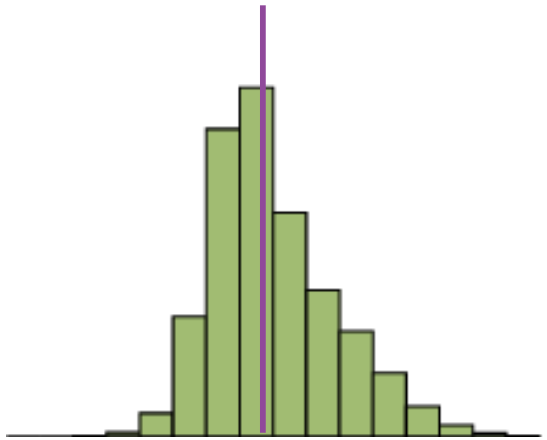

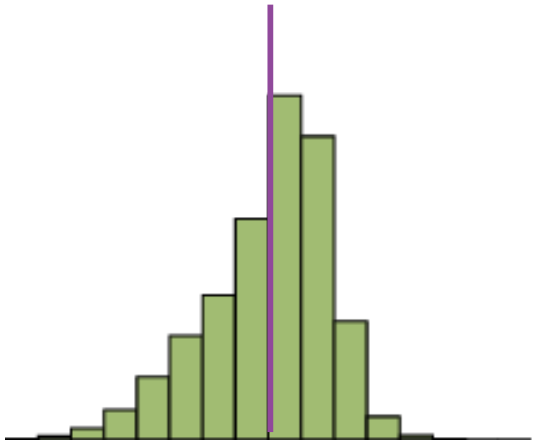

Mean = 0.0
Variance = 1.0

Demo

Feature normalization and standardization using Scikit-Learn

Model Persistence

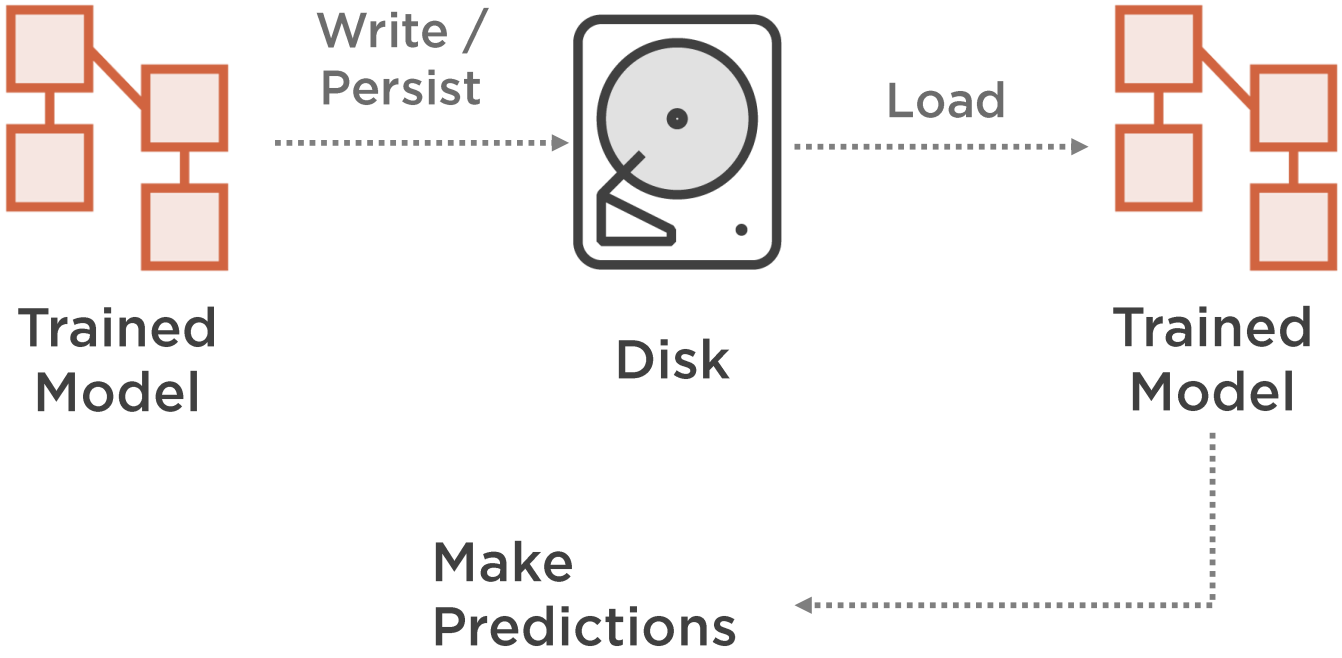

- No retraining required
- Share trained model
- Can be used to create machine learning APIs

Demo

Model persistence using Pickle

Machine Learning API Development

REST API

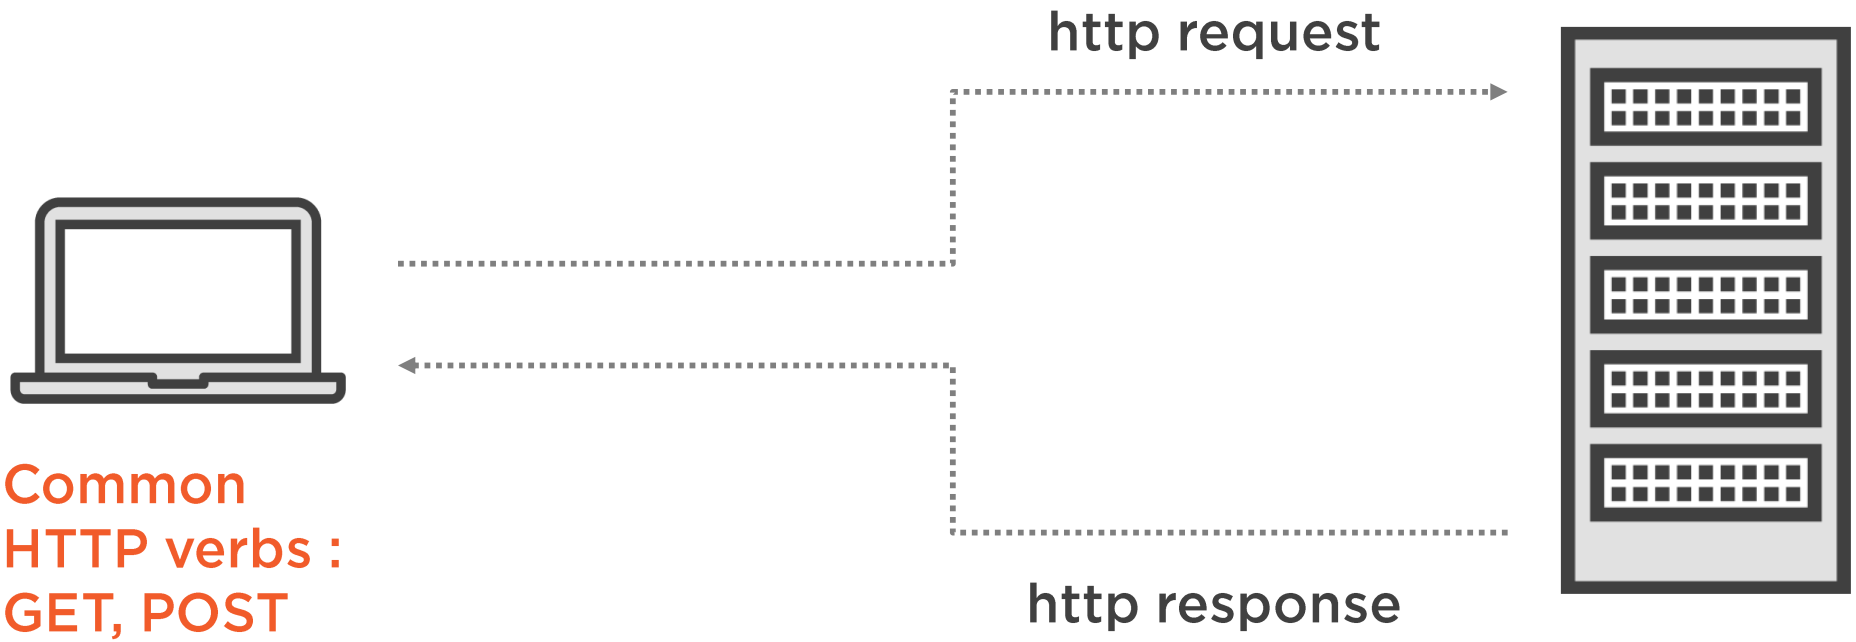

Common
HTTP verbs :
GET, POST

Requests

Machine Learning API Development

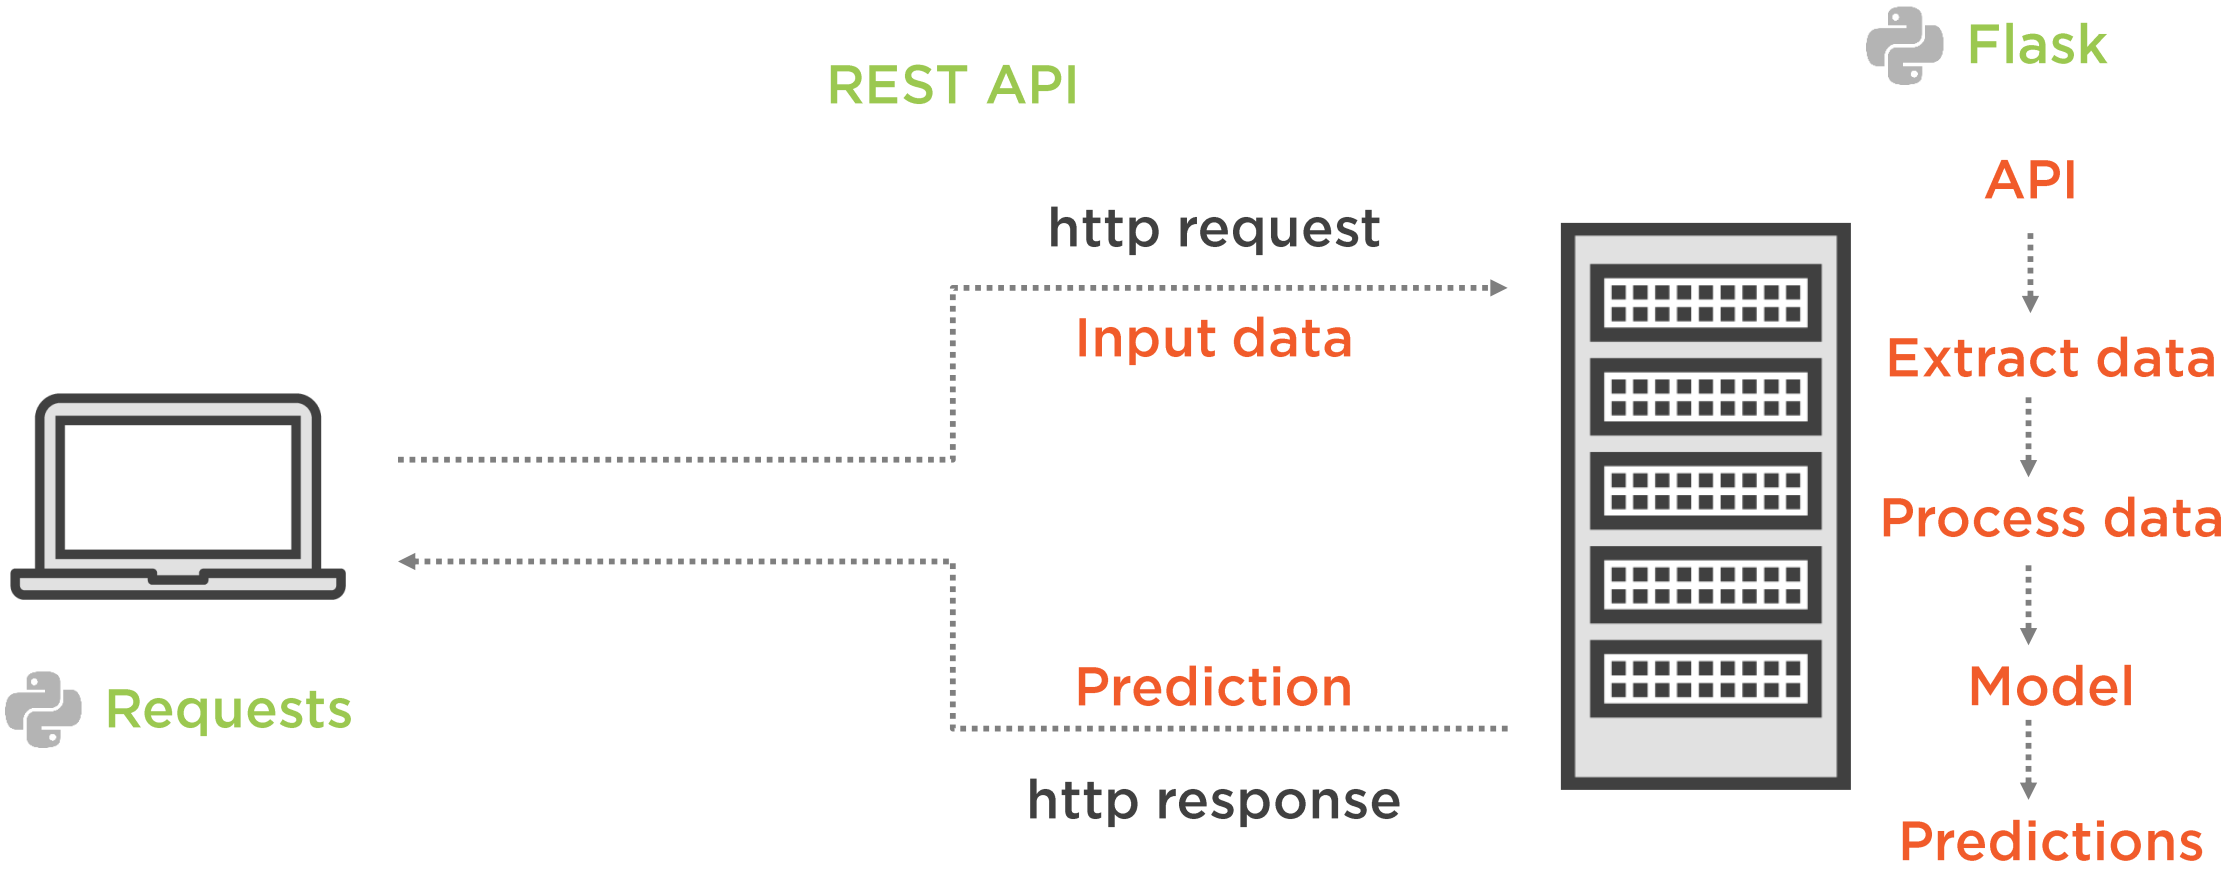

Demo

Hello world API using Flask

Demo

Machine Learning API using Flask

Demo

Committing changes to git

Summary

Model tuning

- Underfitting vs overfitting
- Hyperparameter tuning

Feature normalization

Model persistence

Machine learning API development

Where to Go from Here?

Datasets

Machine learning algorithms

Pipelines, API

Community

Competitions